

Cloud-Native Online Bookstore on Kubernetes

DSAA4040 Engineering-Oriented Project | Group 5

1. Problem Statement

1.1 Objective

The goal of this project is to design, implement, and deploy a **production-ready, cloud-native online bookstore system** on Kubernetes. The system demonstrates modern cloud engineering practices including containerization, microservice-style decomposition, declarative infrastructure management, automated testing, comprehensive observability, and security hardening.

1.2 Motivation

Cloud-native applications are the industry standard for deploying scalable and resilient systems. This project serves as a hands-on exploration of how a real-world e-commerce backend is architected using:

- **Containers** for consistent packaging across environments
- **Kubernetes** for orchestration, scaling, and self-healing
- **Observability** for understanding system behavior under load
- **Security** for protecting user data and business logic
- **Caching** for performance optimization under high load

By building a complete bookstore application from frontend to database, we gain practical experience in the full lifecycle of a cloud-native system: design → implement → test → deploy → monitor → optimize.

1.3 Scope


The implemented system supports the following core features:


Feature	Description
Browse & Search Books	Paginated book catalog with title/author search, Redis caching
Shopping Cart	Add, update, remove items with session isolation and IDOR protection
Order Placement	ACID transaction: order creation, stock deduction, cart cleanup
Payment Mock	Simulated payment gateway with order state transitions
User Authentication	JWT-based registration and login with role-based access
Admin Dashboard	Protected metrics view for administrative users
Auto-scaling	HPA scales backend (2→5) and frontend (1→3) by CPU utilization
Monitoring & Alerting	Prometheus metrics, Grafana dashboards, Alertmanager rules

2. System Design

2.1 Architecture Overview

The system follows a **three-tier architecture** deployed on a single-node Kubernetes cluster (Minikube):





2.2 Technology Stack

Layer	Technology	Rationale
Frontend	Vue 3.4 + Vite + Tailwind CSS + nginx	Modern reactive UI, fast build, minimal container footprint
Backend	Python 3.11 + Flask + Gunicorn	Lightweight, flexible, extensive ecosystem
Database	PostgreSQL 15 + Alembic	Reliable relational store with schema migration support
Cache	Redis 7 (with memory fallback)	Distributed caching for multi-replica consistency; graceful degradation
Auth	PyJWT + Werkzeug	Industry-standard token-based authentication
Orchestration	Kubernetes + Kustomize	Declarative, version-controlled infrastructure
Observability	Prometheus + Grafana + Alertmanager	Metrics collection, visualization, and proactive alerting
CI/CD	GitHub Actions 8-stage pipeline	Automated lint, test, build, security scanning, and deployment validation

2.3 Kubernetes Resources

The application is deployed using **Kustomize** with a base + overlay structure:

Resource	Purpose
Deployment	Manages backend, frontend, PostgreSQL, and Redis pods
Service	ClusterIP for backend/DB/Redis; NodePort for frontend
Ingress	Routes <code>/</code> to frontend and <code>/api/*</code> to backend
HPA	Auto-scales backend (2→5) and frontend (1→3) by CPU
PDB	Ensures minimum 1 pod available during disruptions
NetworkPolicy	Zero-trust mesh: ingress→frontend→backend→db/redis only
ConfigMap/Secret	Non-sensitive config and DB credentials
PVC	Persistent storage for PostgreSQL
ServiceMonitor	Prometheus scrape configuration
PrometheusRule	Alertmanager alert definitions (5 critical alerts)
Grafana Dashboard	Pre-configured 6-panel monitoring dashboard
InitContainer	Runs <code>alembic upgrade head</code> before app starts

2.4 Data Flow

A typical user request flows through the system as follows:

1. **Browser** sends request to `bookstore.local`
2. **Ingress Controller** routes to frontend or backend based on path
3. **Frontend Pod** (nginx) serves Vue SPA or proxies `/api/` to backend
4. **Backend Pod** (Flask) processes request:
5. Check **Redis cache** for book listings (fallback to memory if Redis unavailable)
6. Query PostgreSQL via connection pool if cache miss
7. Return JSON response with `X-Request-ID` header
8. **Prometheus** scrapes `/metrics` endpoint every 15s
9. **Alertmanager** evaluates rules and triggers alerts if thresholds exceeded

3. Implementation

3.1 Backend Architecture

The backend is organized using **Flask Blueprints**, with each domain as a separate module:

```
src/backend/
├── app.py                # Application factory with rate limiting, security headers
├── routes/
│   ├── books.py         # Book catalog (cached, paginated)
│   ├── cart.py          # Shopping cart with IDOR protection
│   ├── orders.py        # Order creation with ACID transactions
│   ├── auth.py          # JWT registration/login
│   ├── payments.py      # Payment state machine
│   ├── admin.py         # Admin-only endpoints
│   └── probes.py        # Health checks and Prometheus metrics (now with DB pool metrics)
├── utils/
│   ├── db.py            # Connection pool + transaction context manager + pool metrics
│   ├── cache.py         # Redis cache with memory fallback
│   ├── auth.py          # JWT generation and verification
│   ├── metrics.py       # Prometheus-style counters (added pool metrics)
│   └── response.py      # JSON serialization with Decimal support
├── schemas.py           # Pydantic v2 validation models
└── alembic/             # Database migration scripts (4 versions)
```

Key Design Decisions:

- 1. Connection Pool Management:** A `ThreadedConnectionPool` (min=1, max=20) is initialized at startup. A custom `db_transaction()` context manager wraps multi-step operations (order creation) in explicit `BEGIN/COMMIT/ROLLBACK` blocks.
- 2. Real-time Pool Metrics:** The `_update_pool_metrics()` function tracks `db_pool_used_connections` and `db_pool_free_connections` on every `getconn() / putconn()`, enabling real-time capacity monitoring via `/metrics` and Grafana.
- 3. Distributed Cache with Fallback:** The `utils/cache.py` module attempts to connect to Redis first. If Redis is unavailable (e.g., in local development), it seamlessly falls back to an in-memory dictionary. This ensures the application works out-of-the-box without Redis.
- 4. Lazy `psycopg2` Imports:** All route files import `RealDictCursor` from `utils.db` rather than directly from `psycopg2`, allowing the application to start and serve fallback mode even when PostgreSQL is not installed.

5. **Pydantic Error Beautification:** The `format_validation_errors()` helper converts raw Pydantic jargon (`Input should be a valid string`) into user-friendly messages (`must be a valid string`), preventing internal field names from leaking to clients.

3.2 Frontend Architecture

The frontend is a **Vue 3 Single-Page Application** built with Vite and served by nginx:

```
src/frontend/src/
├── main.js           # App bootstrap
├── router/index.js   # Route guards (auth, admin)
├── api/client.js     # Axios instance with JWT interceptor
├── store.js         # Reactive cart count
├── components/
│   ├── BookList.vue # Browse + search
│   ├── BookCard.vue # Individual book display
│   ├── CartView.vue # Shopping cart management
│   ├── OrdersView.vue # Order history
│   ├── LoginView.vue # Authentication
│   └── AdminView.vue # Metrics display
```

Key Features: - **JWT Interceptor:** Automatically injects `Authorization: Bearer` header and handles 401 globally - **Route Guards:** `/profile` and `/admin` require authentication; `/login` redirects authenticated users - **Session Management:** `session_id` persisted in `localStorage` for cart isolation

Build Configuration: - `.dockerignore`: Excludes `node_modules/` and `e2e/` from Docker build context - `Dockerfile`: Uses `npm ci` (not `npm install`) for reproducible builds - `vite.config.js`: Development proxy forwards `/api/` to `localhost:8000` - `playwright.config.js`: E2E test configuration with Vite dev server auto-start

3.3 Database Design

The PostgreSQL schema evolved through 4 Alembic migrations:

Table	Key Columns	Purpose
books	id, title, author, isbn, price, stock_quantity	Product catalog
carts	id, session_id	

Table	Key Columns	Purpose
		Shopping cart sessions
cart_items	id , cart_id , book_id , quantity	Items in cart
orders	id , session_id , status , status_history (JSONB)	Order records
order_items	id , order_id , book_id , quantity	Order line items
users	id , username , email , password_hash , is_admin	Authentication

Performance Indexes (Migration 004): - `idx_books_title_author` — accelerates search queries - `idx_books_isbn` — unique lookup - `idx_orders_session_created` — order history queries - `idx_cart_items_cart_book` — cart operations - `idx_order_items_order` — order detail retrieval

3.4 Security Implementations

Feature	Implementation
Authentication	PyJWT HS256 tokens, 24h expiry, Werkzeug password hashing
Authorization	@jwt_required decorator; admin routes check is_admin flag
Rate Limiting	Flask-Limiter with Redis backend (memory fallback). Register: 5/min, Login: 10/min, Default: 200/min
IDOR Prevention	Cart PUT/DELETE verifies ownership via JOIN carts ON session_id
Security Headers	X-Content-Type-Options , X-Frame-Options , HSTS , Referrer-Policy
Container Security	Non-root user (runAsUser: 1000), dropped capabilities, allowPrivilegeEscalation: false
Network Isolation	NetworkPolicy restricts traffic: ingress→frontend→backend→db/redis only

Feature	Implementation
Vulnerability Scanning	Trivy scans on every CI build; results uploaded to GitHub Security tab
Input Validation	Pydantic v2 strict mode; user-friendly error messages without internal field leakage

4. Deployment

4.1 Local Development

The entire system can be deployed locally on Minikube with a single command:

```
minikube start --driver=docker --cpus=2 --memory=4096
minikube addons enable ingress metrics-server
make deploy-tag VERSION=v3.0.3
```

The `Makefile` orchestrates the full workflow: 1. `docker build` — builds backend and frontend images 2. `minikube image load` — transfers images into the Minikube node 3. `kustomize edit set image` — updates overlay with new tags 4. `kubectl apply -k` — deploys all K8s resources 5. `kubectl apply -k k8s/monitoring/` — deploys Grafana dashboard

4.2 Access Methods

Method	URL	Notes
Ingress	<code>http://bookstore.local</code>	Requires <code>/etc/hosts</code> entry
NodePort	<code>http://\$(minikube ip):30080</code>	Direct frontend access
Port-forward	<code>kubectl port-forward</code>	For debugging Prometheus/Grafana

4.3 Kubernetes-Specific Features

InitContainer for Migrations: Every backend pod starts with an `initContainer` that runs `alembic upgrade head`, ensuring the database schema is current before the main container accepts traffic.

Three-Tier Health Probes:

Probe	Endpoint	Purpose	Max Wait
Startup	/startup	Wait for DB pool + migrations	60s
Liveness	/healthz	Detect deadlock; restart if failed	30s
Readiness	/ready	Remove from service endpoints if DB down	15s

Graceful Shutdown: Flask catches `SIGTERM / SIGINT`, closes the DB connection pool, and exits cleanly. Gunicorn runs with `--graceful-timeout 30 --timeout 120` to drain in-flight requests.

Topology Spread Constraints: Backend pods are distributed across nodes using `topologySpreadConstraints` (`maxSkew: 1`, `topologyKey: kubernetes.io/hostname`, `whenUnsatisfiable: ScheduleAnyway`), preventing all replicas from scheduling on the same node.

5. Evaluation & Testing

5.1 Test Strategy

The project employs a **multi-layer testing pyramid**:

Layer	Count	Tools	Description
Unit Tests	75	pytest + mock	Fast, isolated, no DB required
Integration Tests	6	pytest + testcontainers	Real PostgreSQL via Docker
Backend E2E	1	pytest + Playwright (API)	Full HTTP flow against live server
Frontend Unit	25	Vitest + Vue Test Utils	Vue components + API client
Frontend E2E	3	Playwright (Browser)	Browser automation: login → cart → order
Load Tests	3	k6	

Layer	Count	Tools	Description
			HPA trigger, shopping flow, pagination

Total: 115 automated tests

5.2 Test Coverage

Backend coverage is measured at **90%** using `pytest-cov` :

Module	Coverage	Notes
<code>routes/</code>	88%	All endpoints have happy-path and error tests
<code>utils/</code>	92%	Cache, auth, metrics, response helpers
<code>schemas.py</code>	95%	Pydantic validation models

5.3 Performance Results

k6 Load Test (50 Virtual Users, 40 seconds):

Metric	Result
p(95) Latency	1.76ms
Error Rate	< 10%
Throughput	~1,200 req/s

HPA Scaling Test: Under sustained high load (150 VUs), the backend scales from 2 → 5 replicas within 90 seconds. CPU utilization triggers the scaling decision; metrics-server polls every 15 seconds.

5.4 CI/CD Pipeline

GitHub Actions runs an 8-stage pipeline on every push/PR:

1. **Lint** — Hadolint (Dockerfiles), kubeconform (K8s manifests)
2. **Backend Unit Tests** — pytest with coverage report (required)
3. **Backend Integration Tests** — testcontainers + Playwright API (optional, `continue-on-error`)

4. **Frontend Unit Tests** — Vitest component tests + `vite build` (required)
5. **Frontend E2E Tests** — Playwright browser automation (optional, `continue-on-error`)
6. **Build** — Docker Buildx with GitHub Actions cache
7. **Security Scan** — Trivy SARIF → GitHub Security tab
8. **Deploy Check** — Kustomize build validation

Key CI Design Decisions: - Integration and E2E tests run with `continue-on-error: true` so they don't block the pipeline when external dependencies (Docker socket, Playwright browsers) are unavailable - Required jobs (`test-backend` , `test-frontend`) must pass before `build` and `security-scan` - Artifact upload uses `if-no-files-found: warn` to prevent failures when optional test reports don't exist

6. Limitations and Future Work

6.1 Current Limitations

Limitation	Impact	Mitigation
Single-node Minikube	No true HA; node failure = total outage	Use multi-node K3s or cloud K8s for production
Redis single replica	Cache layer is not HA	Deploy Redis Sentinel or Redis Cluster
No HTTPS	Traffic is unencrypted	Add cert-manager + Let's Encrypt in production
No message queue	Order processing is synchronous	Introduce Kafka/RabbitMQ for async tasks
No distributed tracing	X-Request-ID is manual	Integrate Jaeger/Tempo for OpenTelemetry
PostgreSQL single instance	No read replicas or failover	Deploy Patroni or CloudNativePG operator

6.2 Future Work

1. **Service Mesh:** Integrate Istio/Linkerd for automatic mTLS, traffic mirroring, and canary deployments
 2. **GitOps:** Use ArgoCD to automatically sync K8s state from Git, enabling PR-based previews
 3. **Vault Integration:** Replace K8s Secrets with HashiCorp Vault for dynamic credential issuance
 4. **Chaos Engineering:** Regularly inject pod failures and network partitions using Litmus/Chaos Mesh
 5. **FinOps:** Implement Kubecost for namespace-level cost tracking and resource optimization
-

7. Conclusion

This project demonstrates a **complete, production-oriented cloud-native application** built from the ground up. Key achievements include:

- **Architecture:** A clean three-tier design with proper separation of concerns
- **Kubernetes:** Comprehensive use of Deployments, Services, Ingress, HPA, PDB, NetworkPolicy, ConfigMaps, and topology spread constraints
- **Caching:** Redis distributed cache with graceful memory fallback, ensuring multi-replica consistency
- **Observability:** Prometheus metrics (including real-time DB pool monitoring), Grafana dashboards, Alertmanager rules, and X-Request-ID tracing
- **Security:** JWT authentication, rate limiting, IDOR fixes, security headers, non-root containers, and vulnerability scanning
- **Reliability:** Graceful shutdown, three-tier health probes, ACID transactions, and initContainer migrations
- **Quality:** 115 automated tests with 90% backend coverage, multi-layer testing pyramid, and 8-stage CI/CD pipeline

The system successfully deploys on Minikube with a single command (`make deploy-tag`) and provides a solid foundation for production deployment on any Kubernetes cluster.

8. References

1. Kubernetes Documentation. <https://kubernetes.io/docs/>
 2. Flask Documentation. <https://flask.palletsprojects.com/>
 3. Vue.js 3 Guide. <https://vuejs.org/guide/>
 4. Prometheus Monitoring. <https://prometheus.io/docs/>
 5. Alembic Migrations. <https://alembic.sqlalchemy.org/>
 6. Minikube Quick Start. <https://minikube.sigs.k8s.io/docs/start/>
 7. Cloud Native Computing Foundation. <https://www.cncf.io/>
-

Report updated on May 20, 2026